



LCI Advanced Workshop 2025: Slurm High-Availability & Federation

Josh Burks
HPC Systems Analyst
Arizona State University Research Computing
josh.burks@asu.edu

*This document is a result of work by volunteer LCI instructors and is licensed under CC BY-NC 4.0
(<https://creativecommons.org/licenses/by-nc/4.0/>)*

Learning Goals

- Learn what Slurm High-Availability is and isn't, and when it makes sense to use
- Understand how to configure and verify a multi-controller HA setup
- Appreciate the trade-offs that led our site to abandon Slurm HA
- Learn what Slurm Federation is, and isn't, and when it makes sense to use it
- Understand basic federation setup (sacctmgr commands, cluster roles) and federated job submission
- Apply best practices for testing, monitoring, and maintaining both HA and federation setups

HA Architecture

- Slurm has always supported HA. The code for it exists in the 1.0 branch of slurm, with references in the changelog as far back as SLURM 0.2.12.
- Was significantly revised in 18.08 to support more than one backup controllers.
- Slurm HA is a fallback system - one controller is deemed primary, all other controllers are backup

HA Architecture

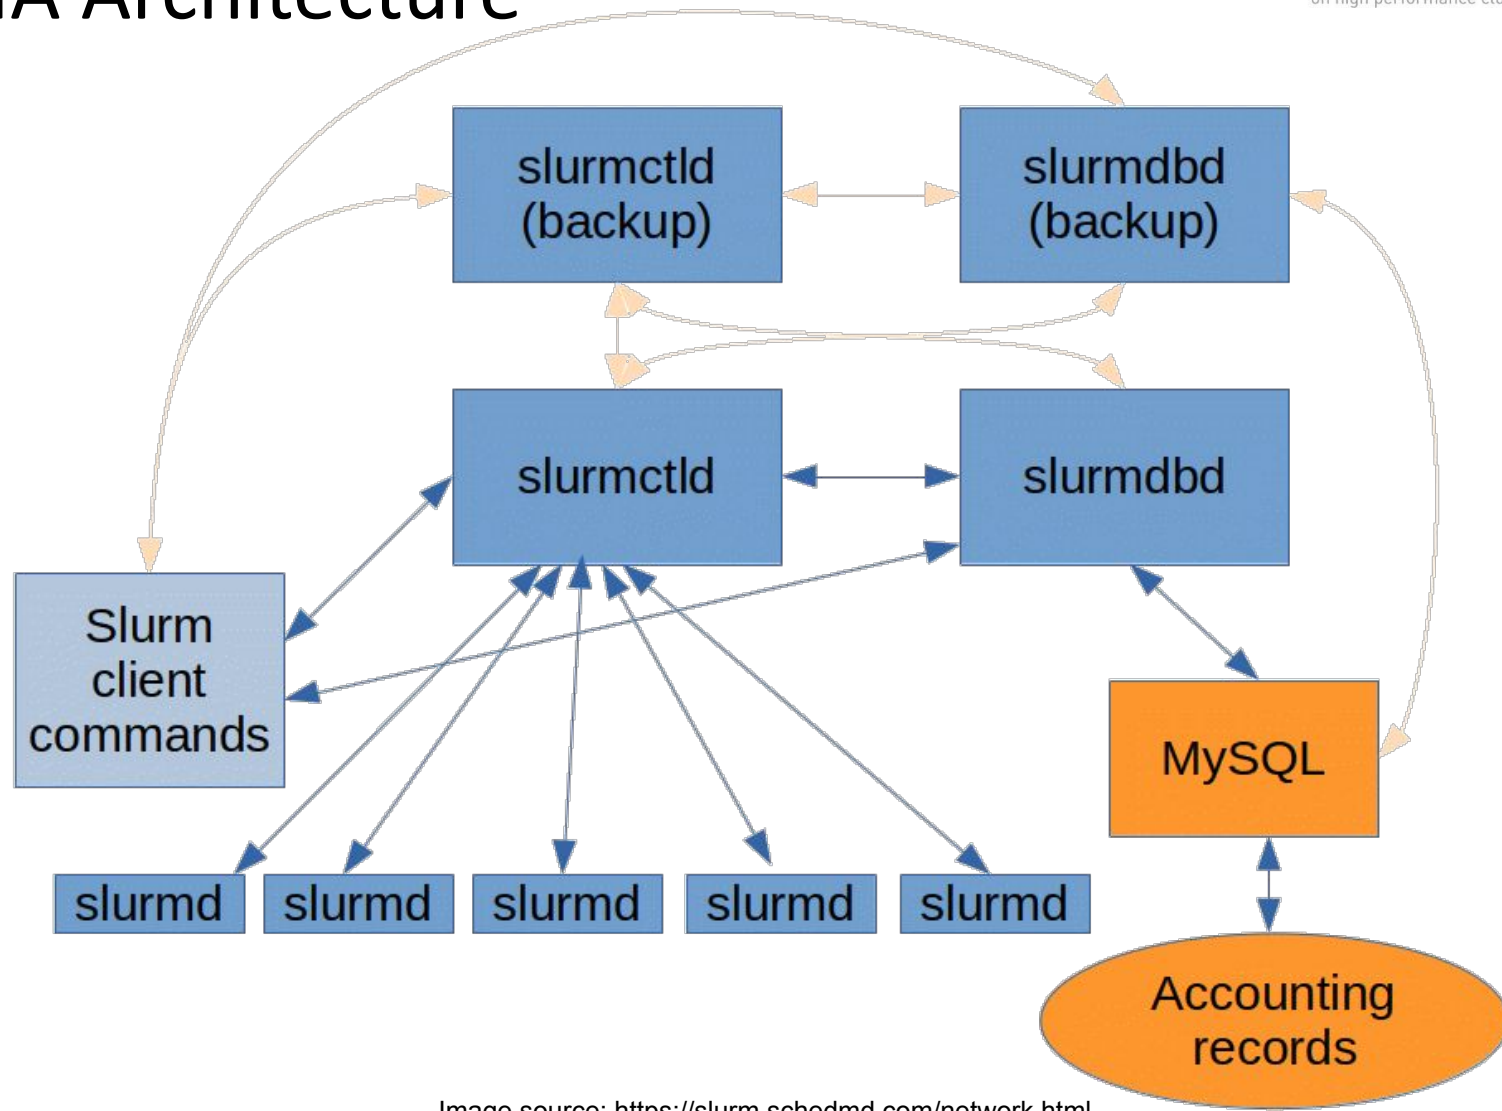


Image source: <https://slurm.schedmd.com/network.html>

HA Architecture

- A shared filesystem is required on all slurmctld hosts (StateSaveLocation)
- This is where slurms stores its “state” - queued, running and recently completed jobs, node state
- If a controller comes up without access to the state information, queued and running jobs will be cancelled.
- This shared file system should NOT be an NFS mounted filesystem
 - Slurm does constant read/write operations to this location, and NFS is not up to the task

HA Configuration

```
#/etc/slurm/slurm.conf
```

```
SlurmctldHost=slurm01(10.10.30.1) # Primary controller
```

```
#/etc/slurm/slurm.conf
```

```
SlurmctldHost=slurm01(10.10.30.1) # Primary controller
```

```
SlurmctldHost=slurm02(10.10.30.2) # Secondary controller
```

```
# scontrol reconfig
```

HA Configuration

```
#/etc/slurm/slurm.conf
```

```
SlurmctldHost=slurm01(10.10.30.1)  # Primary controller
```

```
SlurmctldHost=slurm02(10.10.30.2)  # Secondary controller
```

```
SlurmctldPrimaryOnProg=/usr/local/bin/slurmctld_takeover.sh
```

```
SlurmctldPrimaryOffProg=/usr/local/bin/slurmctld_release.sh
```

```
SlurmctldTimeout=60
```

```
SlurmctldParameters=no_backup_scheduling
```

```
# scontrol reconfig
```



ASU Research Computing



ASU[®] Research Computing

HA Best Practices

- Perform failover drills (bring primary down, verify backup takes over)
- Use a small value for `slurmctldTimeout` to ensure quick failover in the event of an unexpected shutdown
- Maintain tight time synchronization across controllers and nodes
- Use reliable, low-latency shared storage for StateSaveLocation
 - SAN, DRDB, LINSTOR, but not NFS

HA Alternatives

- Slurmctld in a VM
 - Less prone to hardware failures
 - Migratable
- Slurmctld in a Container
 - Use Kubernetes for HA
- Federation
 - Other cluster's slurmctld can take over some job scheduling


What Is Federation?

- Federation = loose coordination of multiple Slurm clusters.
- Allows job and job step ID consistency across clusters.
 - Bits 0-25: Local Job ID
 - Bits 26-31: Cluster Origin ID
- Enables job forwarding between clusters when local resources are insufficient.
- Federated clusters retain full autonomy; no shared state or database.
- No global scheduler - each cluster operates independently but shares limited information.

Multi-Cluster vs Federation

- Multi-Cluster
 - Each cluster has its own slurmctld(s) and may use a shared or separate slurmdbds
 - If slurmdbds are separate, they must be linked with with AccountingStorageExternalHost= in slurm.conf
 - Users have separate accounts for each cluster
 - Users who have an account on both clusters can submit jobs to either cluster with the -M flag
- Federated Clusters
 - Each cluster has its own slurmctld(s) but must use **shared** slurmdbd(s)
 - Joined via the shared database
 - Unified accounting structure

Multi-Cluster



Federation




Diagram source: <https://slurm.schedmd.com/network.html>

Federation Architecture

- Each cluster runs its own SlurmDBD and controller (slurmctld).
- Federation relies on:
 - slurmctld communication
 - job submission tools (e.g., sbatch, srun)
- Uses slurmdbd federation support for job/account tracking.
- Consistent UID/GID & account naming is critical.
- Technically a shared file system is not required, but it is recommended
- Requires bidirectional network visibility between slurmctld instances.

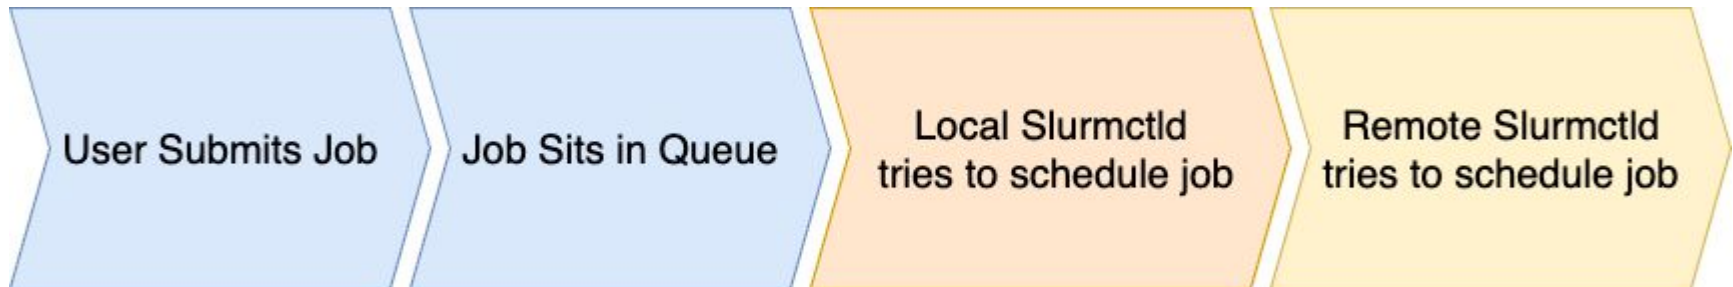
Federation Configuration

- Enable federation with:

```
sacctmgr create federation name=myfed clusters=cluster1,cluster2
```


Job Lifecycle in Federation

- User submits job
- Job sits in queue
- Local slurmctld tries to schedule it
- If resources are unavailable, it considers forwarding (if configured)
- Remote slurmctld accepts and schedules job



```

Every 1 hr: state = clusterw.      cluster: Sat Oct 4 20:08:13 2008
Cluster: cluster1
PARTITION Aged:  TIMELEFT:  NODES:  STATE: NODELIST:
default      up    infinite      2  all-compute[01-02]
Cluster: cluster2
PARTITION Aged:  TIMELEFT:  NODES:  STATE: NODELIST:
default      up    infinite      1  all-fab-compute[01-02]

```

Federation Best Practices & Limits

- Best Practices:
 - Ensure consistent UID/GID, accounts, and partitions across clusters.
 - Use identical job submission wrappers or front-ends to avoid surprises.
 - Test forwarding with low-priority jobs before production usage.
 - Monitor with tools like sacct, squeue with --clusters.

- Limitations:
 - No **global** job prioritization or fair-share enforcement.
 - No shared scheduling view across clusters.
 - Forwarded jobs can behave differently based on remote cluster policy.
 - Debugging issues is more complex across cluster boundaries.
 - Job arrays only run on the origin cluster.

Q & A