



LCI Advanced Workshop 2025: Resource Limits

Josh Burks
HPC Systems Analyst
Arizona State University Research Computing
josh.burks@asu.edu

*This document is a result of work by volunteer LCI instructors and is licensed under CC BY-NC 4.0
(<https://creativecommons.org/licenses/by-nc/4.0/>)*

Outline

- Explore some of the different mechanisms for controlling the different resources different users can access
- Explore how condo models interact with resource limits
- Discuss the limits we can set to
 - Maintain fairness between research and classroom workloads
 - Enforce allocations that match funding or ownership
 - Keep utilization high while maintaining predictable access
- Explore how job priority is assigned
- Explore how nodes are selected for jobs

Resource Enforcement

| Target | Role in Resource Limits | Example of Controls |
|-----------|--|---|
| User | Per-User Limits | MaxTRES, MaxJobs |
| Account | Groups users for accounting and fairshare | GrpTRES, MaxJobsPerAccount |
| QoS | Defines policy, priority, and quotas | MaxTRES, MaxWall, Preemption, PriorityWeightQOS |
| Partition | Groups nodes for scheduling and access control | MaxNodes, MaxTime, AllowAccounts, PriorityTier |

Association, QoS, and Partitions

- Accounts are “billing” accounts
- Associations map users to accounts to clusters and are the primary vehicle for long-term fairshare calculator and per-group limits
- QoS’s create policy around scheduling priority and overriding limits, including the impact to fairshare with per-user/account TRES limits, fairshare decay, UsageFactors
- Partitions create groups of nodes with access policies and scheduling queues with basic job-level limits

Account Limits

- Accounts define who can use resources and how much they can consume.
- Limits cascade through hierarchy:
 - User -> Account -> Cluster -> Global
- Departmental chargeback or usage tracking
- Restricting total cores, GPUs, or jobs per group
- Enforcing per-user quotas inside shared accounts

```
sacctmgr modify account research set \  
GrpTRES=cpu=2048 MaxJobs=100
```

Association Limits

- Association is the mapping of User define who can use resources and how much they can consume.
- Limits cascade through hierarchy:
 - User -> Account -> Cluster -> Global

```
sacctmgr modify user where name=josh account=research set \  
MaxJobs=10 MaxSubmitJobs=20
```

Departmental Hierarchy

Parent-Child Hierarchical Limits

```
sacctmgr add account research parent=root GrpTRES=cpu=8000
sacctmgr add account chem parent=research GrpTRES=cpu=2000
sacctmgr add account pi_joe parent=chem GrpTRES=cpu=512
sacctmgr add user josh account=pi_joe MaxTRESPerJob=cpu=32
```

Departmental Hierarchy

Parent-Child Hierarchical Limits

```
# sacctmgr list association tree \
format=Account,User,GrpTRES,MaxTRES
```

| Account | User | GrpTRES | MaxTRES |
|--------------|-------|-----------|---------|
| ----- | ----- | ----- | ----- |
| root | | cpu=20000 | |
| root | root | | |
| research | | cpu=8000 | |
| college_chem | | cpu=2000 | |
| pi_jones | | cpu=512 | cpu=512 |
| pi_jones | josh | | cpu=32 |

QoS Fundamentals

- QoS provides flexible controls: priority weight, max jobs, max walltime, max TRES, preemption flags, partition override flags
- PriorityWeightQOS multiplies into job priority to favor or deprioritize groups (Normalized value)
- QoS can set MaxTRESPerUser, MaxTRESPerJob, MaxSubmitJobs and enforce quotas
- QoS may be configured to preempt lower-QoS jobs

QoS Resource Limits

- Classroom policy: small MaxWall (2h), limited CPUs , fewer GPUs, higher QoS priority
- Research policy: larger MaxWall, larger MaxTRES limits, higher-priority QoS, access to private nodes.

```
sacctmgr add qos research Priority=100 \  
MaxWall=72:00:00
```

```
sacctmgr add qos classroom Priority=200 \  
MaxWall=02:00:00 MaxSubmitJobs=2 \
```

QoS Resource Limits:

- Default research QoS

```
sacctmgr create qos name=research
```

QoS Resource Limits:

- Longer run time

```
sacctmgr create qos name=long \  
    MaxJobsPerUser=2 \  
    MaxWallDurationPerJob=14-00:00:00 \  
    MaxTRESPerUser=cpu=512 \  
    Flags=PartitionTimeLimit  
  
sacctmgr modify user where user=jeburks2 \  
    account=pi_joe qos+=long
```

QoS Resource Limits:

- Debugging QoS

```
sacctmgr create qos name=debug \  
  Priority=20000000 \  
  MaxJobsPerUser=2 \  
  UsageFactor=0.5 \  
  MaxWallDurationPerJob=15 \  
  Flags=DenyOnLimit
```

QoS Resource Limits:

- Idle Cycles

```
sacctmgr create qos idlecycles \  
UsageFactor=0.001 \  
Flags=DenyOnLimit,NoReserve
```

QoS Resource Limits:

- Condo Node Owners

```
sacctmgr create qos pi_joe \  
  Priority=500000 \  
  Preempt=idlecycles \  
  UsageFactor=0 \  
  Flags=PartitionTimeLimit,DenyOnLimit,OverPartQOS
```

Assumes: PreemptType=preempt/qos

QoS Resource Limits:

- Classroom Limited Use

```
sacctmgr create qos name=class \  
Priority=200 \  
MaxJobsPerUser=2 \  
MaxWallDurationPerJob=1440 \  
MaxTRESPerUser=gres/gpu=1,cpu=32,mem=120G
```


Flexible Time-Based Limits

- Sometimes fixed limits are too broad
- Maybe limiting classroom users to 1 GPU is too strict
- Maybe classroom users need 4 GPUs, but for a shorter duration
- Flexible time-based limits allow users longer wall times when they use less resources
- Use MaxTRESRunMinsPU and MaxTRESRunMinsPerAccount

Maybe we want to classroom academic courses users run at most jobs combining to 16 hours (960 minutes) of GPU time

1 GPU for 16 Hours / 2 GPUs for 12 Hours / 3 GPUs for 8 Hours / 4 GPUs for 4 Hours

QoS Resource Limits:

- Classroom Limited Use

```
sacctmgr create qos name=class \  
Priority=200 \  
MaxJobsPerUser=2 \  
MaxWallDurationPerJob=1440 \  
MaxTRESRunMinsPU=gres/gpu=960 ,cpu=768
```

QoS in the Scheduling Formula

- Job priority in Slurm is calculated from multiple weighted factors:
- $\text{Priority} = \text{Age} + \text{Fairshare} + \text{JobSize} + \text{Partition} + \text{QOS} + \text{Site}$
- QoS priority contributes through PriorityWeightQOS in slurm.conf.
- Example:
 - $\text{PriorityWeightQOS} = 1000000$
 - Higher weight = stronger impact of QoS on total job priority.
 - Use `sprio --long` to inspect the breakdown for queued jobs.
 - Combine with PriorityTier in partitions to fine-tune scheduling order.

Analyzing Priority

```
# sprio --long
```

| JOBID | PARTITION | USER | PRIORITY | AGE | FAIRSHARE | PARTITION | QOSNAME | QOS |
|----------|-----------|-------|-----------|----------|-----------|-----------|---------|-----------|
| 33553665 | public | userA | 10000009 | 10000000 | 10 | 0 | public | 0 |
| 34374167 | highmem | userB | 10000011 | 10000000 | 11 | 0 | public | 0 |
| 34374170 | highmem | userB | 10000011 | 10000000 | 11 | 0 | public | 0 |
| 34398154 | public | userA | 10000009 | 10000000 | 10 | 0 | public | 0 |
| 34556382 | highmem | userC | 10000228 | 10000000 | 228 | 0 | public | 0 |
| 34556395 | highmem | userC | 10000228 | 10000000 | 228 | 0 | public | 0 |
| 34556399 | highmem | userC | 10000228 | 10000000 | 228 | 0 | public | 0 |
| 34834438 | general | userD | 219843814 | 10000000 | 99843815 | 0 | grp_c | 110000000 |
| 34852751 | general | userE | 168124584 | 10000000 | 97624585 | 0 | debug | 60500000 |
| 34854466 | public | userF | 10632112 | 10000000 | 632113 | 0 | public | 0 |
| 34864821 | highmem | userG | 10261876 | 10000000 | 261876 | 0 | public | 0 |
| 34871673 | highmem | user | 10000000 | 10000000 | 0 | 0 | public | 0 |

Analyzing Priority

```
Set DebugFlags to NO_CONF_HASH,Priority
debug: slurmctld log levels: stderr=debug logfile=debug syslog=quiet
Set debug level to 'debug'
priority/multifactor: _get_fairshare_priority: PRIO: Fairshare priority of job 5 for user root in acct root is
2**(-1.000000/1.000000) = 0.500000
priority/multifactor: _get_priority_internal: Weighted Age priority is 0.000000 * 10000000 = 0.00
priority/multifactor: _get_priority_internal: Weighted Assoc priority is 0.000000 * 0 = 0.00
priority/multifactor: _get_priority_internal: Weighted Fairshare priority is 0.500000 * 100000000 = 50000000.00
priority/multifactor: _get_priority_internal: Weighted JobSize priority is 0.000000 * 0 = 0.00
priority/multifactor: _get_priority_internal: Weighted Partition priority is 0.000000 * 0 = 0.00
priority/multifactor: _get_priority_internal: Weighted QOS priority is 0.000000 * 110000000 = 0.00
priority/multifactor: _get_priority_internal: Site priority is 0
priority/multifactor: _get_priority_internal: Job 5 priority: 0 + 0 + 0.00 + 50000000.00 + 0.00 + 0.00 + 0.00 + 0 -
0 = 50000000.00
PRIO: BillingWeight: JobId=5 is either new or it was resized
PRIO: BillingWeight: JobId=5 using "CPU=1.0,Mem=.25G,gres/gpu=3.0" from partition general
PRIO: BillingWeight: JobId=5 SUM(TRES) = 1.500000
sched: _slurm_rpc_allocate_resources JobId=5 NodeList=lci-compute-XX-2 usec=673
```

Preemption with QoS

- QoS can define who gets to reclaim resources when the system is full
- Preemption modes determine what happens to lower-priority jobs
 - REQUEUE / CANCEL / SUSPEND
- Preemption rules are controlled by PreemptMode and PreemptType:

PreemptType=preempt/qos

PreemptMode=REQUEUE # REQUEUE / CANCEL / SUSPEND

Condo Model

- PI or research group contributes hardware ('buys in') to a shared cluster
- Nodes integrated into the global system for scheduling and accounting
- Owners expect priority/guaranteed access, opportunistic use of other resources
- Admins want high utilization, enforceable policies, fairness across groups

ASU Condo Implementation

- Approach:
 - One 'condo' partition for all condo owners.
 - Each group has a QoS with a priority boost
 - Condo-owners nodes are targeted with JSP
 - Non-owner jobs may run on idle condo nodes with the idlecycles QoS with preemption
- Benefits:
 - Simple configuration (1 partition + per-group QoS).
 - Owners rarely wait on their nodes.
 - Idle time minimized.
- Drawbacks
 - In order to enforce node restrictions via JSP, editing of submitted jobs is disabled

Other Condo Strategies

- Dedicated Partitions or QoS (per group):
 - Strict node ownership, owners only
 - Pros: Clear ownership, owners own a specific piece of hardware, simple job submission for users
 - Cons: Many partitions to manage
- Reservation-Based Model:
 - Guarantees “equivalent hardware” via reservations.
 - Pros: flexible, better utilization.
 - Cons: owners may dislike non-dedicated nodes.
- Hybrid Approaches:
 - Nodes in owner-only and shared partitions.
 - QoS ensures owners get preference
 - Balances efficiency and ownership

Preemption in Condo Models

- Why Preemption?
 - Ensures owners regain condo resources quickly.
 - Protects guarantees while allowing opportunistic use.
- Strategies:
 - PreemptMode=REQUEUE - requeue outsider jobs.
 - PreemptMode=CANCEL - cancel outsider jobs outright.
 - PreemptMode=SUSPEND - pause jobs, resume later.
- QoS preemption rules - owners' QoS can preempt others.
- Trade Offs:
 - REQUEUE: fair, but can frustrate users.
 - CANCEL: harsh, instant access for owners.
 - SUSPEND: jobs stay in memory
 - QoS layering: flexible but complex.

Partition Fundamentals

- Partitions are logical groups of nodes that define where jobs can run
- Each partition can have unique limits, scheduling behavior, and access policies
- Common parameters:
 - Nodes= list of nodes in the partition
 - Default=YES/NO jobs submitted without a -p flag will go there
 - PriorityTier= affects scheduling order relative to other partitions
 - MaxTime=, MaxNodes=, MaxCPUsPerUser= resource limits
 - AllowAccounts=, AllowQos= restrict who can submit
- Typical use cases:
 - Separate CPU, GPU, and high-memory nodes
 - Isolate public/condo nodes

Partition Fundamentals

- Partitions can also inherit **some** of the resource limits of a QoS
- Only inherits limits that can already apply to partitions
 - Does not assign jobs to the QoS
 - Does not give the job any priority characteristics of the QoS
 - Does not give the job any preemption characteristics of the QoS
 - Does not inherit accounting limits (ie, time-based limits)

Partition Fundamentals

- Partitions are like “queues” in PBS/Torque - jobs are grouped and prioritized within each partition.
- Thinking of them as queues helps conceptualize how Slurm schedules jobs.
- The main scheduler processes partitions in order of PriorityTier.

Partition Fundamentals

- When it hits a high-priority job that can't start (e.g., needs 64 GPUs unavailable), it stops evaluating lower-priority jobs in that same partition for that pass
- The backfill scheduler, which runs afterward, can still start smaller or lower-priority jobs if they fit without delaying higher-priority ones
- This behavior means a single large job can block smaller jobs in the same partition.
- Creating dedicated partitions for resource-intensive hardware (e.g., GPUs, large-memory nodes) prevents these bottlenecks and improves overall throughput.

Node Weights

- Assuming all other factors are equal and several idle nodes can satisfy the requested resources - how does slurm pick which node to run on?
- Node “weights” in slurm.conf influence the order in which nodes are allocated
- Higher Weight = lower scheduling preference
- Useful for favoring newer, faster, or high-bandwidth nodes

```
nodeName=compute[001-199] Weight=1000 # Slow, Old Nodes
```

```
nodeName=compute[200-299] Weight=1 # Fast, New Nodes
```

Node Weights

- Try it on your lab VM:
- Submit a job, and notice how you land on node 1

```
[root@lci-head-XX-1 ~]# salloc -p general,
```

```
salloc: Granted job allocation 3
```

```
salloc: Nodes lci-compute-XX-1 are ready for job
```


Node Weights

- Try it on your lab VM:
- Update your nodes to assign different weights, and notice how you land on node 2

```
scontrol update nodename=lci-compute-XX-1 weight=1000
```

```
[root@lci-head-XX-1 ~]# salloc -p general
```

```
salloc: Granted job allocation 4
```

```
salloc: Waiting for resource configuration
```

```
salloc: Nodes lci-compute-XX-2 are ready for job
```

Node Weights

- To make this permanent, assign weight in /etc/slurm/slurm.conf to assign your nodes different weights

```
NodeName=lci-compute-XX-[1-2] \
```

```
Weight=1 \
```

```
CPUs=2 \
```

```
Boards=1 \
```

```
SocketsPerBoard=2 \
```

```
CoresPerSocket=1 \
```

```
ThreadsPerCore=1 \
```

```
RealMemory=7500
```

Node Weights

- To make this permanent, assign weight in /etc/slurm/slurm.conf to assign your nodes different weights

```
NodeName=lci-compute-XX-1 \  
Weight=1000 \  
CPUs=2 \  
Boards=1 \  
SocketsPerBoard=2 \  
CoresPerSocket=1 \  
ThreadsPerCore=1 \  
RealMemory=7500
```

```
NodeName=lci-compute-XX-2 \  
Weight=1 \  
CPUs=2 \  
Boards=1 \  
SocketsPerBoard=2 \  
CoresPerSocket=1 \  
ThreadsPerCore=1 \  
RealMemory=7500
```

Q & A