



LCI Advanced Workshop 2025: **Security in Slurm**

Josh Burks
HPC Systems Analyst
Arizona State University Research Computing
josh.burks@asu.edu

*This document is a result of work by volunteer LCI instructors and is licensed under CC BY-NC 4.0
(<https://creativecommons.org/licenses/by-nc/4.0/>)*

Learning Goals

- Understand Slurm's built-in security mechanisms
- Recognize weaknesses in traditional HPC security models
- Explore built in options for hardening Slurm clusters
- Explore advanced tools (SPANK, containers, UserBasedFirewall)
- Identify ongoing risks beyond Slurm's scope

Historical HPC Security Model

- “Hard exterior, soft interior”
 - External security: VPNs, bastion hosts, MFA
 - Internal security: basic POSIX permissions, SELinux often disabled
- Users don’t have full free roam, but inside is loosely controlled
- Onion model: layers exist, but mostly at the perimeter
- Malicious insiders (or compromised accounts) remain a threat

Unencrypted HPC Communications

- MPI and RDMA traffic are unencrypted
- Often bypass host firewalls (RDMA direct access)
- Mitigation:
 - VLAN separation for management vs. compute
 - UserBasedFirewall for cross-user traffic control

Slurm Security Layers

- Transport Security (TLS + certmgr)
- Job & Node Isolation (MCS, containers, /proc controls)
- Authentication & Authorization (PAM, Slurm Adopt, AdminLevels)
- Data Privacy & Audit Control (PrivateData)

PrivateData

- Restrict visibility of accounts, jobs, usage, reservations
- Protects sensitive data in shared accounting systems

#slurm.conf

```
PrivateData=jobs,accounts,events,jobs,nodes,partitions,rese  
rvations,usage,users
```

#slurmdbd.conf

```
PrivateData=accounts,events,jobs,reservations,usage,users
```

Node Sharing & /proc Security

- Node sharing risks → job snooping
- /proc isolation with hidepid=1
 - Users may not access any /proc/<pid>/ directories but their own, protected against local eavesdroppers.
 - Reduces exposure of other users' processes

```
# /etc/fstab
```

```
proc /proc proc hidepid=1 0 0
```

Job Container Tmpfs

- Isolated /tmp file systems
- Configurable job-level filesystem isolation
- If userA writes sensitive data to /tmp, normally userB could see this
- By bind-mounting /tmp to a private per-job mount, userB can only see userB's /tmp

Job Container Tmpfs

```
#slurm.conf
```

```
JobContainerType=job_container/tmpfs
```

```
PrologFlags=Contain
```

```
#job_container.conf
```

```
AutoBasePath=true
```

```
BasePath=/mnt/job_tmp # Real location where data lives
```

```
Dirs=/tmp,/dev/shm # Paths bind-mounted under the basepath
```

```
Shared=true # Required when using autofs
```

MCS (Multi-Category Security)

- Tenant isolation via categories
- Ensures jobs with different categories never share nodes
- Useful where data may be seen by members of a group, but not across groups on a shared system
 - `mcs/none` disables MCS labels and functionality.
 - `mcs/account` MCS labels equal the job's `--account`
 - `mcs/group` MCS labels equal to the job's user group
 - `mcs/user` MCS labels equal to the username of the job's `--uid`
 - `mcs/label` MCS labels are arbitrary strings

PAM Slurm Adopt

- PAM plugin that prevents users from sshing into nodes on which they don't have a running job
- Prevents rogue or orphaned processes
- The user's connection is "adopted" into the extern step cgroup of the job so that they cannot exceed cgroup limits
- All processes created by the user and the user's connection are killed when the job ends

PAM Slurm Adopt

```
#/etc/pam.d/sshd
```

```
account      required      pam_nologin.so
```

```
account      include      password-auth
```

```
...
```

```
-account      required      pam_slurm_adopt.so \  
               action_generic_failure=deny \  
               action_adopt_failure=deny
```

PAM Slurm Adopt + SELinux

- PAM Slurm Adopt requires a custom SELinux module to work in an environment where SELinux is enabled

```
module pam_slurm_adopt 1.0;  
require {  
    ...  
}  
allow sshd_t  
...
```

TLS + certmgr + certgen

- New plugins in Slurm 25.05
- Encrypts all Slurm RPC traffic
- Protects job submissions, scheduling, accounting
- Uses s2n (signal-to-noise) github.com/aws/s2n-tls
 - dnf install s2n-tls-devel

TLS + certmgr + certgen

```
#slurm.conf
```

```
TLSType=tls/s2n
```

```
TLSParameters=ca_cert_file=/etc/pki/slurm_ca.pem
```

```
CertmgrType=certmgr/script
```

```
CertgenType=certgen/script
```

```
#slurmdbd.conf
```

```
TLSType=tls/s2n
```

```
TLSParameters=ca_cert_file=/etc/pki/slurm_ca.pem
```

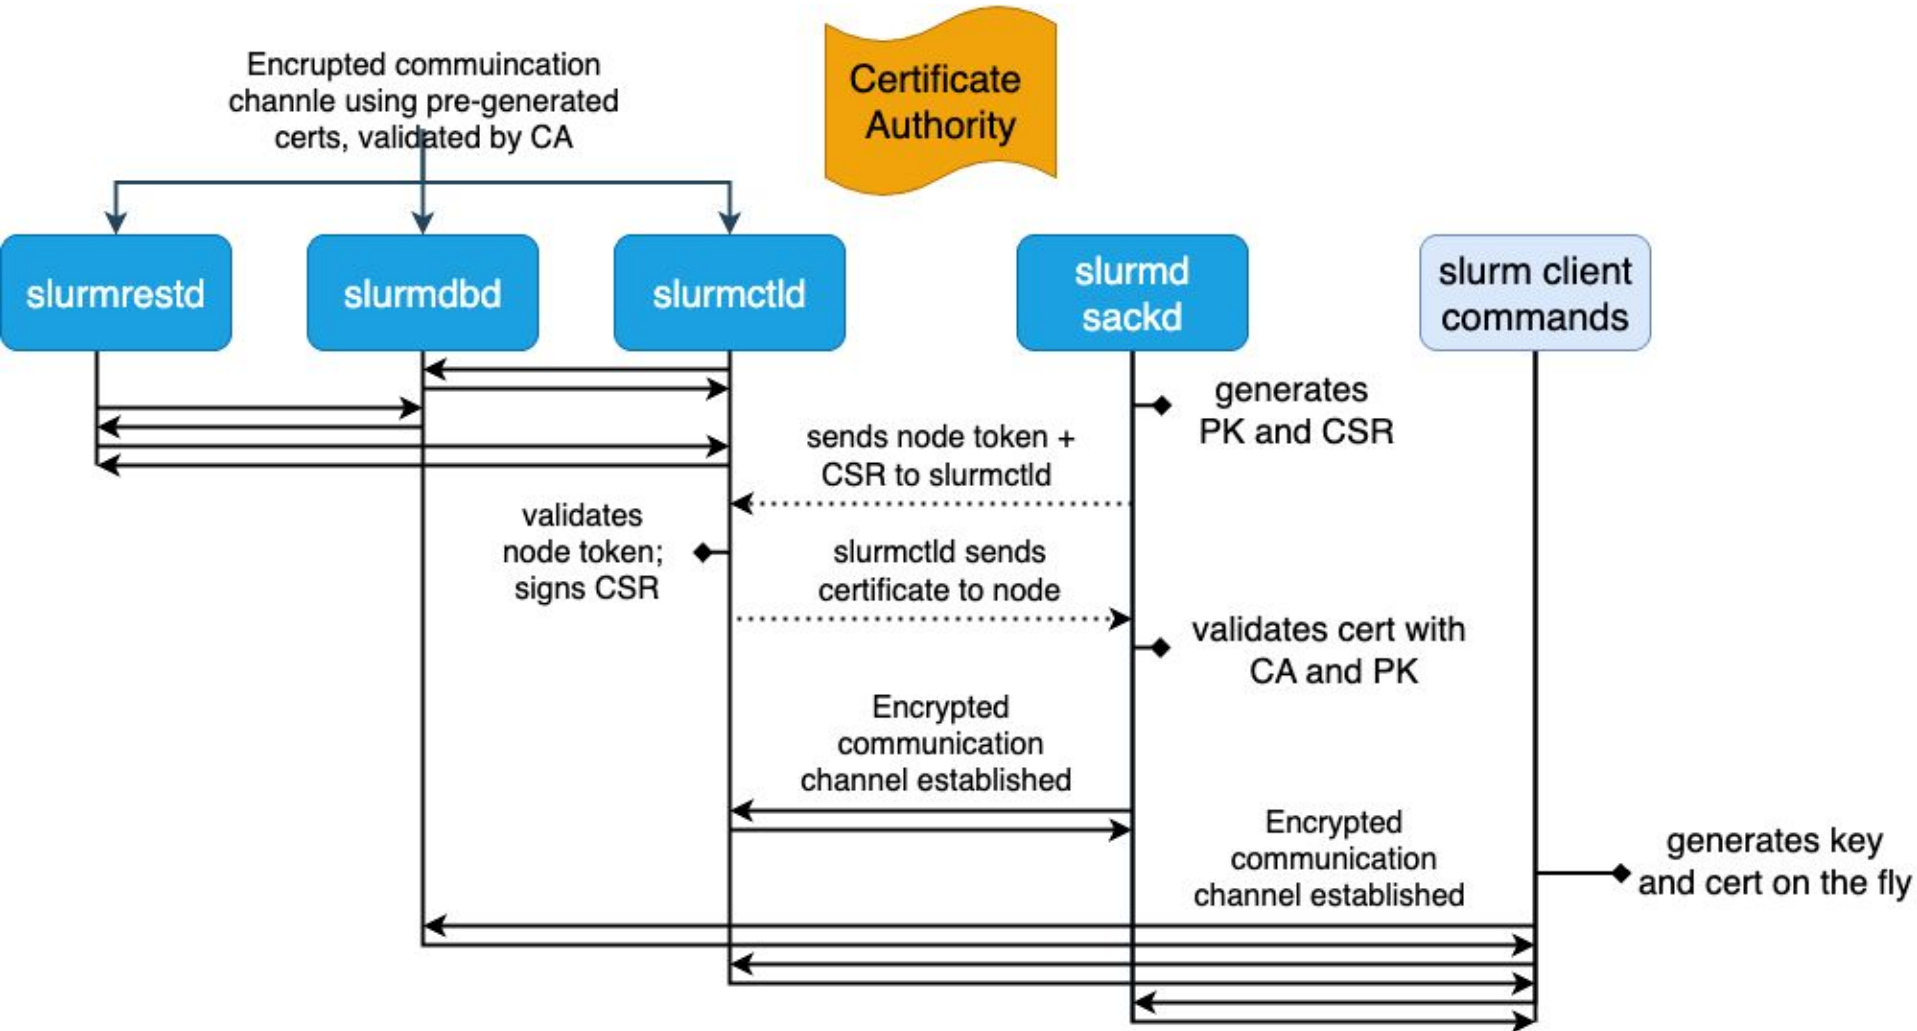
TLS + certmgr + certgen

- certmgr plugin manages certificates for slurmd/sackd
- certgen plugin generates key/cert pairs on the fly for client commands
- slurmctld, slurmdbd, and slurmrestd all have unique certificates
- slurmd boot with a pre-shared token
- slurmd generates a private key
- slurmd sends token+certificate signing request to slurmctld
- slurmctld validates token is owned, and returns a certificate
- sackd follows the same process

Note:

- all certificates are signed by a common certificate authority
- slurmd must be started with `-ca-cert-file`

TLS + certmgr + certgen



Beyond Slurm: Network Security

- MPI & RDMA unencrypted
- Infiniband fabric controls: partition keys, SR-IOV
- VLAN separation (mgmt vs compute)
- UserBasedFirewall for user-to-user traffic control
 - github.com/mit-llsc/UserBasedFirewall

Firewalling & Ports

- Limit which hosts can connect to Slurm daemons
- Use firewalls to restrict RPC traffic
- Control SlurmctldPort and SlurmdPort
- Limit ephemeral ports srun can use
- CommunicationParameters=block_null_hash (new in 21.08.8)
- ,NoCtldInAddrAny,NoInAddrAny

Firewalling & Ports

- Limit which hosts can connect to Slurm daemons
- Use firewalls to restrict RPC traffic
- Control SlurmctlPort and SlurmdPort
- Limit ephemeral ports srun can use
- Block null hash (new in 21.08.8)

```
#slurm.conf
```

```
SlurmctlPort=6817
```

```
SlurmdPort=6818
```

```
SrunPortRange=61000-62000
```

```
CommunicationParameters=block_null_hash,NoCtldInAddrAny,NoInAddrAny
```

Prolog/Epilog for Security

- Use them to sanitize hosts between jobs
- Reset GPUs, wipe sensitive data
- Purge old mounts
- Ensure the job directory is private

SPANK Plugins

- Extend Slurm job launch with site-specific security policies
- Examples: enforce job isolation, restrict env vars, audit logging
 - github.com/BYUHPC/oodproxy

DOS & Abuse Prevention

- Rate-limit per-user RPC
- Limit number of jobs a user can submit
- Prevent sbatch fork-bombs

#slurm.conf

SlurmctldParameters=rl_enable

Staying Secure with Slurm

- Watch SchedMD advisories for CVEs
- Paid customers get early advisories
- Patch early, test often
- Security is an ongoing process

Putting It All Together

- TLS + certmgr -> encrypted daemon comms
- MCS + containers + hidepid -> isolate jobs and processes
- PAM + SPANK -> enforce user/job controls
- Prolog/Epilog sanitization -> clean nodes securely
- PrivateData + auditing -> protect accounting
- Network firewalls + VLANs + fabric keys -> protect backplane
- Fairshare enforcement and rate limiting RPC -> prevent abuse

Q & A