



# LCI Advanced Workshop 2025: Slurm Scheduler Tuning: Balancing HTC & Large Jobs

Josh Burks  
HPC Systems Analyst  
Arizona State University Research Computing  
[josh.burks@asu.edu](mailto:josh.burks@asu.edu)

*This document is a result of work by volunteer LCI instructors and is licensed under CC BY-NC 4.0  
(<https://creativecommons.org/licenses/by-nc/4.0/>)*

# Learning Goals

- Understand the differences between HTC and large-job workloads and their impact on scheduling
- Learn how to identify and mitigate large-job starvation in mixed environments
- Explore scheduler parameter tuning for different cluster sizes
- Compare the roles of the main scheduler vs backfill scheduler
- Explore best practices for partition design and public/private node management
- Gain strategies for iterative tuning and monitoring performance impact

# Resource Allocation Models in Slurm

## Whole Node Allocation (SelectType=select/linear)

- Job gets entire node regardless of requested resources
- Simple, predictable, no fragmentation
- Wasteful for small jobs

# Resource Allocation Models in Slurm

## Consumable Resources (SelectType=select/cons\_tres)

- Allocates CPUs, GPUs, memory individually
- SelectTypeParameters: CR\_CPU, CR\_Core, CR\_Socket, CR\_Memory
- Maximum resource utilization but increased complexity

## Scheduling Considerations:

- Job sizing and resource requests
- Node fragmentation
- Memory enforcement (MemEnforceLimit)
- GPU/accelerator placement
- NUMA awareness

# Scheduler Types in Slurm

## **SchedulerType=sched/buildin**


- FIFO scheduler initiates jobs in priority order.
- If any job in the partition can not be scheduled, no lower priority job in that partition will be scheduled.
- exception is made for jobs that can not run due to partition constraints (e.g. the time limit) or down/drained nodes.

## **SchedulerType=sched/backfill**

- Backfill scheduling module to augment the default FIFO scheduling
- Effectiveness dependent upon users specifying job time limits

# Main vs Backfill Scheduler Roles

- Main scheduler: determines job start order based on priority
- Backfill scheduler: fills idle resources without delaying higher-priority jobs



# Main Scheduler Loop

- Wake up every SchedulerTimeSlice seconds (default: 5)
- Load job queue sorted by priority
- For each job in priority order:
  - Check partition limits (MaxNodes, MaxCPUsPerUser, etc.)
  - Check association/QOS limits
  - Check job dependencies
  - Attempt resource allocation
  - If allocation succeeds: start job or make reservation
  - If allocation fails: continue to next job
- Stop when max\_sched\_time exceeded OR bf\_max\_job\_test jobs evaluated
- Sleep until next cycle

# Backfill Scheduler Loop

- Triggered every `bf_interval` seconds (default: 30)
- Build "reservation table" from main scheduler decisions
- Calculate earliest start times for all pending jobs
- For each job in queue (not priority order):
  - If job can start now without violating reservations: start it
  - Track jobs started (`bf_max_job_start` limit)
  - Track jobs tested (`bf_max_job_test` limit)
- Stop when limits hit or queue exhausted



# The Value of Backfill Scheduling

## Without Backfill

- Strict priority queue processing
- Resources idle while waiting for large jobs
- Poor cluster utilization
- Users frustrated by idle nodes they can't access

## With Backfill

- - Smaller jobs fill gaps between large jobs
- - Better user experience (shorter wait times for small jobs)
- - Maintains fairshare and priority goals

# Balancing Partitions to Optimize Scheduling

## Overlapping Resources

- Multiple partitions sharing same nodes
- Scheduler evaluates each partition separately
- Duplicate effort checking same resources
- Can cause scheduler performance issues

## Minimize partition overlap where possible

- Use AllowAccounts/AllowGroups instead of separate partitions
- Partition by job characteristics (size, runtime, priority)
- Monitor sdiag for excessive partition evaluation times

# Using sdiag to Diagnose Scheduler Issues

```
# sdiag
```

```
Main schedule statistics
```

```
(microseconds) :
```

```
    Last cycle:    2302
```

```
    Max cycle:    62634
```

```
    Total cycles: 26014
```

```
    Mean cycle:    3323
```

```
    Mean depth cycle: 69
```

```
    Cycles per minute: 57
```

```
    Last queue length: 68
```

```
Main scheduler exit:
```

```
    End of job queue:26012
```

```
    Hit default_queue_depth: 2
```

```
    Hit sched_max_job_start: 0
```

```
    Blocked on licenses: 0
```

```
    Hit max_rpc_cnt: 0
```

```
    Timeout (max_sched_time): 0
```

# Using sdiag to Diagnose Scheduler Issues

```
# sdiag
```

```
Backfilling stats
```

```
Total backfilled jobs: 243060
Total cycles: 454
Last cycle: 609383
Max cycle: 853384
Mean cycle: 294402
Last depth cycle: 130
Last depth cycle (try sched): 130
Depth Mean: 189
Depth Mean (try depth): 188
```

```
Backfill exit
```

```
End of job queue:454
Hit bf_max_job_start: 0
Hit bf_max_job_test: 0
System state changed: 0
Hit table size limit: 0
Timeout (bf_max_time): 0
```

# Tuning Main Scheduler Cycle Timing

- **max\_sched\_time** - Max time per main cycle (default: 2s), prevents scheduler from blocking too long
- **sched\_interval** - How often the main scheduling loop will run, considering all jobs while still honoring the `partition_job_depth` limit
- **default\_queue\_depth** - Number of jobs to consider for scheduling on each event that may result in a job being scheduled. (Default 100)

# Tuning Backfill Scheduler Cycle Timing

## **bf\_interval**

- How often backfill runs (default: 30s)
- Independent of main scheduler cycle
- Reduce for fast-moving HTC workloads
- Increase to reduce overhead on large systems

# Tuning Backfill Scheduler Cycle Timing

## **bf\_window**

- Jobs needing resources beyond window ignored
- Increase for long jobs
- A value at least as long as the highest allowed time limit is generally advisable to prevent job starvation

## **bf\_window\_linear**

- Setting this will create higher resolution of time precision in the backfill scheduler uses linear backoff instead of exponential back off
- Not recommended with more than a few hundred simultaneously executing jobs.

# Tuning Backfill Scheduler Cycle Timing

## **bf\_max\_job\_test**

- Max jobs evaluated per backfill cycle (default: 300)
- Increase for deeper backfill
- Watch cycle times - can make cycles slow

## **bf\_max\_job\_start**

- Max jobs started per backfill cycle (default: 0 = unlimited)
- Limit to prevent backfill from being too aggressive
- Useful for protecting large job opportunities

## **bf\_max\_job\_assoc**

- How many jobs per user associona backfill can check
- Large impact when users submit many jobs



# Tuning for All HTC Workloads

**defer** - Do not attempt to schedule jobs individually at submit time. Can be useful for high-throughput computing.

**defer\_batch** - Like **defer**, but only will defer scheduling for batch jobs.

**pack\_serial\_at\_end** - may reduce resource fragmentation by put serial jobs at the end of the available nodes

**SelectTypeParameters=CR\_LLN** - Schedule resources to jobs on the least loaded nodes

# Workload Mix & Scheduling Pressure

## HTC Workloads Want:

- Fast turnaround (seconds to minutes)
- High throughput (jobs/second)
- Frequent backfill opportunities
- Minimal wait times

## Large Jobs Want:

- Resource guarantees
- Protected start times
- Priority enforcement
- Minimal fragmentation

# Tuning for All HTC

# Fast, frequent cycles

SchedulerTimeSlice=2

SchedulerParameters=max\_sched\_time=1

# Aggressive backfill

SchedulerParameters=bf\_interval=15

SchedulerParameters=bf\_window=120 # 2 hours

SchedulerParameters=bf\_max\_job\_test=1000

SchedulerParameters=bf\_max\_job\_start=0 # unlimited

SchedulerParameters=bf\_continue

PriorityFavorSmall=Yes

# Tuning for All Large Jobs

# Standard cycles, deeper evaluation

SchedulerTimeSlice=5

SchedulerParameters=max\_sched\_time=4

# Conservative backfill

SchedulerParameters=bf\_interval=60

SchedulerParameters=bf\_window=10080 # 7 days

SchedulerParameters=bf\_max\_job\_test=300

SchedulerParameters=bf\_max\_job\_start=5

SchedulerParameters=bf\_reserve\_nodes=1

# Workload Mix & Scheduling Pressure

## The Conflict:

- HTC jobs fill all available resources
- Large jobs can't find contiguous allocations
- Priority system gets subverted by constant backfill
- Users of both types dissatisfied

Solution Requires careful balance of main and backfill tuning, partition design, and policy enforcement

# Tuning for Hybrid Workloads

# Moderate cycle timing

SchedulerTimeSlice=5

SchedulerParameters=max\_sched\_time=2

# Tuned backfill

SchedulerParameters=bf\_interval=30

SchedulerParameters=bf\_window=1440 # 24 hours

SchedulerParameters=bf\_max\_job\_test=500

SchedulerParameters=bf\_max\_job\_assoc=20 # CRITICAL

# Scheduler Tuning at ASU

```
SchedulerParameters      =  
bf_window=43200  
bf_resolution=300  
bf_max_time=120  
bf_max_job_assoc=25  
bf_max_job_test=150000  
bf_interval=30  
max_rpc_cnt=1000  
enable_user_top  
bf_continue  
nohold_on_prolog_fail
```

# Monitoring & Iteration

- Baseline before changes
- Metrics: wait times, utilization, fairshare
- Gradual adjustments
- User feedback loop
- Slurm Exporter:
  - [github.com/vpenso/prometheus-slurm-exporter](https://github.com/vpenso/prometheus-slurm-exporter)
  - [github.com/ubccr/slurm-exporter](https://github.com/ubccr/slurm-exporter) (Uses Slurm API)
  - [gitlab.com/SchedMD/slinky/slurm-exporter](https://gitlab.com/SchedMD/slinky/slurm-exporter) (Official)



# Monitoring & Iteration

[grafana.com/dashboards/4323](https://grafana.com/dashboards/4323)

## Backfill Scheduler Cycles



## Total Backfilled Jobs (since last slurm start) ⓘ



# Runaway Jobs

```
# sacctmgr show runawayjobs
```

**NOTE:** Runaway jobs are jobs that don't exist in the controller but have a start time and no end time in the database

```
[root@sol-slurm01:~]# sacctmgr show runawayjobs
```

NOTE: Runaway jobs are jobs that don't exist in the controller but have a start time and no end time in the database

ID	Name	Partition	Cluster	State	TimeSubmit	TimeStart	TimeEnd
28108882	/home/jeb+	lightwork	sol	PENDING	2025-07-02T00:05:17	Unknown	Unknown
35866256	ood-vscod+	general	sol	PENDING	2025-10-14T09:49:36	Unknown	Unknown
35866258	ood-virtu+	general	sol	PENDING	2025-10-14T09:49:46	Unknown	Unknown
35866283	interacti+	htc	sol	PENDING	2025-10-14T10:39:40	Unknown	Unknown

Would you like to fix these runaway jobs?

(This sets the end time for each job to the latest of the job's start, eligible, and submit times, and sets the state to completed.

Once corrected, this triggers the SlurmDBD to recalculate the usage from before the earliest submit time of all the runaway jobs. **Warning:** This could take a long time and sreport may not return data until the recalculation is complete d.)

(You have 30 seconds to decide)

# Key Takeaways

- Match tuning to workload type and site goals
- Use partitions and limits strategically
- Balance main scheduler and backfill tuning
- Iterate with data and feedback

# Q & A